

STA 235H - Model Selection II: Shrinkage

Fall 2022

McCombs School of Business, UT Austin

Last class



- Started with our **prediction chapter**
 - Bias vs. Variance
 - Validation set approach and Cross-validation
 - How to choose a model for a continuous outcome (RMSE)
 - Stepwise selection

Knowledge check from last week

- 1) Which model is **higher bias**: A complex model or a simpler one?
- 2) Why do we **split our data** into training and testing datasets?
- 3) How do we **compare** models with continuous outcomes?

How forward stepwise selection works: Example from last class

1) Start with a null model (no covariates)

- Your best guess will be the average of the outcome in the training dataset!

2) Test out all models with **one covariate**, and **select the best one**:

- E.g. $\text{logins} \sim \text{female}$, $\text{logins} \sim \text{succession}$, $\text{logins} \sim \text{age}$, ...
- $\text{logins} \sim \text{succession}$ is the best one (according to RMSE)

3) Test out all models with **two covariates**, but that have *succession*!

- E.g. $\text{logins} \sim \text{succession} + \text{female}$, $\text{logins} \sim \text{succession} + \text{age}$, $\text{logins} \sim \text{succession} + \text{city}$, ...

4) You will end up with k possible models (k : total number of predictors).

- Choose the best one, depending on the RMSE.

Today: Continuing our journey

- How to **improve our linear regressions**:
 - Ridge regression
 - Lasso regression
- Look at **binary outcomes**



Honey, I shrunk the coefficients!


What is shrinkage?

- We reviewed the **stepwise procedure**: Subsetting model selection approach.
 - Select k out of p total predictors
- **Shrinkage (a.k.a Regularization)**: Fitting a model with all p predictors, but introducing bias (i.e. shrinking coefficients towards 0) for improvement in variance.
 - Ridge regression
 - Lasso regression

On top of a ridge.


Ridge Regression: An example

- Predict spending based on frequency of visits to a website




Ordinary Least Squares

- In an **OLS**: Minimize sum of squared-errors, i.e. $\min_{\beta} \sum_{i=1}^n (\text{spend}_i - (\beta_0 + \beta_1 \text{freq}_i))^2$




What about fit?

- Does the OLS fit the testing data well?



Ridge Regression

- Let's shrink the coefficients!: Ridge Regression



Ridge Regression: What does it do?

- Ridge regression **introduces bias to reduce variance** in the testing data set.
- In a simple regression (i.e. one regressor/covariate):

$$\min_{\beta} \sum_{i=1}^n \underbrace{(y_i - \beta_0 - x_i \beta_1)^2}_{OLS}$$

Ridge Regression: What does it do?

- Ridge regression **introduces bias to reduce variance** in the testing data set.
- In a simple regression (i.e. one regressor/covariate):

$$\min_{\beta} \sum_{i=1}^n \underbrace{(y_i - \beta_0 - x_i \beta_1)^2}_{OLS} + \underbrace{\lambda \cdot \beta_1^2}_{RidgePenalty}$$

- λ is the **penalty factor** \rightarrow indicates how much we want to shrink the coefficients.

Q1: In general, which model will have smaller β coefficients?


a) A model with a larger λ

b) A model with a smaller λ

**Remember... we care about
accuracy in the testing dataset!**


RMSE on the testing dataset: OLS

$$RMSE = \sqrt{\frac{1}{4} \sum_{i=1}^4 (\text{spend}_i - (132.5 - 16.25 \cdot \text{freq}_i))^2} = 28.36$$



RMSE on the testing dataset: Ridge Regression

$$RMSE = \sqrt{\frac{1}{4} \sum_{i=1}^4 (\text{spend}_i - (119.5 - 11.25 \cdot \text{freq}_i))^2} = 12.13$$



Ridge Regression in general

- For regressions that include **more than one regressor**:

$$\min_{\beta} \underbrace{\sum_{i=1}^n (y_i - \sum_{k=0}^p x_i \beta_k)^2}_{OLS} + \underbrace{\lambda \cdot \sum_{k=1}^p \beta_k^2}_{RidgePenalty}$$

- In our previous example, if we had two regressors, *female* and *freq*:

$$\min_{\beta} \sum_{i=1}^n (\text{spend}_i - \beta_0 - \beta_1 \text{female}_i - \beta_2 \text{freq}_i)^2 + \lambda \cdot (\beta_1^2 + \beta_2^2)$$


- Because the ridge penalty includes the β 's coefficients, **scale matters**:
 - Standardize variables (*you will do that as an option in your code*)

How do we choose λ ?


Cross-validation!

- 1) Choose a grid of λ values
 - The grid you choose will be context dependent (play around with it!)
- 2) Compute cross-validation error (e.g. RMSE) for each
- 3) Choose the smallest one.

λ vs RMSE?



λ vs RMSE? A zoom



How do we do this in R?

```
library(caret)

set.seed(100)

hbo = read.csv("https://raw.githubusercontent.com/...")

lambda_seq = seq(0, 20, length = 500)

ridge = train(logins ~ . - unsubscribe - id,
              data = train.data,
              method = "glmnet",
              preProcess = "scale",
              trControl = trainControl("cv", numfolds = 10),
              tuneGrid = expand.grid(alpha = 0,
                                    lambda = lambda_seq)
)

cv_lambda = data.frame(lambda = ridge$results$lambda_1se,
                       rmse = ridge$results$RMSE)
```

- We will be using the caret package

How do we do this in R?

```
library(caret)

set.seed(100)

hbo = read.csv("https://raw.githubusercontent.com/...")

lambda_seq = seq(0, 20, length = 500)

ridge = train(logins ~ . - unsubscribe - id,
              data = train.data,
              method = "glmnet",
              preProcess = "scale",
              trControl = trainControl("cv", numfolds = 10),
              tuneGrid = expand.grid(alpha = 0,
                                    lambda = lambda_seq)
)
```

- We will be using the `caret` package
- We are doing **cross-validation**, so remember to set a seed!

How do we do this in R?

```
library(caret)

set.seed(100)

hbo = read.csv("https://raw.githubusercontent.com/...")

lambda_seq = seq(0, 20, length = 500)

ridge = train(logins ~ . - unsubscribe - id,
              data = train.data,
              method = "glmnet",
              preProcess = "scale",
              trControl = trainControl("cv", numfolds = 10),
              tuneGrid = expand.grid(alpha = 0,
                                    lambda = lambda_seq)
)

cv_lambda = data.frame(lambda = ridge$results$lambda_1se,
                       rmse = ridge$results$RMSE)
```

- We will be using the `caret` package
- We are doing **cross-validation**, so remember to set a seed!
- You need to create a grid for the λ 's **that will be tested**

How do we do this in R?

```
library(caret)

set.seed(100)

hbo = read.csv("https://raw.githubusercontent.com/...")

lambda_seq = seq(0, 20, length = 500)

ridge = train(logins ~ . - unsubscribe - id,
              data = train.data,
              method = "glmnet",
              preProcess = "scale",
              trControl = trainControl("cv", numfolds = 10),
              tuneGrid = expand.grid(alpha = 0,
                                    lambda = lambda_seq))

cv_lambda = data.frame(lambda = ridge$results$lambda_1se,
                       rmse = ridge$results$RMSE)
```

- We will be using the `caret` package
- We are doing **cross-validation**, so remember to set a seed!
- You need to create a grid for the λ 's **that will be tested**
- The function we will use is `train`: Same as before
 - `method="glmnet"` means that it will run an elastic net.
 - `alpha=0` means is a **ridge regression**
 - `lambda = lambda_seq` is not necessary (you can provide your own grid)

How do we do this in R?

```
library(caret)

set.seed(100)

hbo = read.csv("https://raw.githubusercontent.com/...")

lambda_seq = seq(0, 20, length = 500)

ridge = train(logins ~ . - unsubscribe - id,
              data = train.data,
              method = "glmnet",
              preProcess = "scale",
              trControl = trainControl("cv", numfolds = 10),
              tuneGrid = expand.grid(alpha = 0,
                                    lambda = lambda_seq)
)

cv_lambda = data.frame(lambda = ridge$results$lambda,
                       rmse = ridge$results$RMSE)
```

- We will be using the `caret` package
- We are doing **cross-validation**, so remember to set a seed!
- You need to create a grid for the λ 's **that will be tested**
- The function we will use is `train`: Same as before
- Important objects in CV:
 - `results$lambda`: Vector of λ that was tested
 - `results$RMSE`: RMSE for each λ
 - `bestTune$lambda`: λ that minimizes the error term.

How do we do this in R?

OLS regression:

```
lm1 = lm(logins ~ succession + city,  
          data = train.data)
```

```
coef(lm1)
```

```
## (Intercept)  succession      city  
##      7.035888    -6.306371    2.570454
```

```
rmse(lm1, test.data)
```

```
## [1] 2.089868
```

Ridge regression:

```
coef(ridge$finalModel, ridge$bestTune$lambda)
```

```
## 5 x 1 sparse Matrix of class "dgCMatrix"  
##                                     s1  
## (Intercept)  6.564243424  
## female      0.002726465  
## city        0.824387472  
## age         0.046468790  
## succession  -2.639308962
```

```
rmse(ridge, test.data)
```

```
## [1] 2.097452
```

Throwing a lasso

Lasso regression

- Very similar to ridge regression, except it **changes the penalty term**:

$$\min_{\beta} \underbrace{\sum_{i=1}^n (y_i - \sum_{k=0}^p x_i \beta_k)^2}_{OLS} + \underbrace{\lambda \cdot \sum_{k=1}^p |\beta_k|}_{LassoPenalty}$$

- In our previous example:

$$\min_{\beta} \sum_{i=1}^n (\text{spend}_i - \beta_0 - \beta_1 \text{female}_i - \beta_2 \text{freq}_i)^2 + \lambda \cdot (|\beta_1| + |\beta_2|)$$

- Lasso regression is also called l_1 regularization:

$$||\beta||_1 = \sum_{k=1}^p |\beta_k|$$

Ridge vs Lasso

Ridge

Final model will have p coefficients

Usually better with multicollinearity

Lasso

Can set coefficients = 0

Improves interpretability of model

Can be used for model selection

And how do we do Lasso in R?

```
library(caret)
set.seed(100)

hbo = read.csv("https://raw.githubusercontent.com/...")

lambda_seq = seq(0, 20, length = 500)

lasso = train(logins ~ . - unsubscribe - id, data = hbo,
              method = "glmnet",
              preProcess = "scale",
              trControl = trainControl("cv", number = 10),
              tuneGrid = expand.grid(alpha = 1,
                                    lambda = lambda_seq)
)

cv_lambda = data.frame(lambda = lasso$results$lambda_1se,
                       rmse = lasso$results$RMSE)
```

Exactly the same!

- ... But change $\alpha=1$!!

And how do we do Lasso in R?

Ridge regression:

```
coef(ridge$finalModel, ridge$bestTune$lambda)
```

```
## 5 x 1 sparse Matrix of class "dgCMatrix"
##               s1
## (Intercept)  6.564243424
## female      0.002726465
## city        0.824387472
## age         0.046468790
## succession  -2.639308962
```

```
rmse(ridge, test.data)
```

```
## [1] 2.097452
```

Lasso regression:

```
coef(lasso$finalModel, lasso$bestTune$lambda)
```

```
## 5 x 1 sparse Matrix of class "dgCMatrix"
##               s1
## (Intercept)  6.84122778
## female      .
## city        0.87982819
## age         0.03099797
## succession  -2.83492585
```

```
rmse(lasso, test.data)
```

```
## [1] 2.09171
```

A note on binary outcomes

- If we are predicting **binary outcomes**, RMSE would not be an appropriate measure anymore!
 - We will use **accuracy instead**: The proportion (%) of correctly classified observations.
- For example:

```
set.seed(100)

lasso = train(factor(unsubscribe) ~ . - id, data = train.data,
               method = "glmnet", preProcess = "scale",
               trControl = trainControl("cv", number = 10),
               tuneGrid = expand.grid(alpha = 1, lambda = lambda_seq))

pred.values = lasso %>% predict(test.data)

mean(pred.values == test.data$unsubscribe)
```

```
## [1] 0.736
```

A note on binary outcomes

- If we are predicting **binary outcomes**, RMSE would not be an appropriate measure anymore!
 - We will use **accuracy instead**: The proportion (%) of correctly classified observations.
- For example:

```
set.seed(100)

lasso = train(factor(unsubscribe) ~ . - id, data = train.data,
              method = "glmnet", preProcess = "scale",
              trControl = trainControl("cv", number = 10),
              tuneGrid = expand.grid(alpha = 1, lambda = lambda_seq))

pred.values = lasso %>% predict(test.data)

mean(pred.values == test.data$unsubscribe)
```

```
## [1] 0.736
```

A note on binary outcomes

- If we are predicting **binary outcomes**, RMSE would not be an appropriate measure anymore!
 - We will use **accuracy instead**: The proportion (%) of correctly classified observations.
- For example:

```
set.seed(100)

lasso = train(factor(unsubscribe) ~ . - id, data = train.data,
              method = "glmnet", preProcess = "scale",
              trControl = trainControl("cv", number = 10),
              tuneGrid = expand.grid(alpha = 1, lambda = lambda_seq))

pred.values = lasso %>% predict(test.data)

mean(pred.values == test.data$unsubscribe)
```

```
## [1] 0.736
```

Your Turn

Main takeaway points

- You can **shrink coefficients** to introduce bias and decrease variance.
- Ridge and Lasso regression are **similar**:
 - Lasso can be used for model selection.
- Importance of understanding **how to estimate the penalty coefficient**.



References

- James, G. et al. (2021). "Introduction to Statistical Learning with Applications in R". *Springer. Chapter 6*.
- STDHA. (2018). "Penalized Regression Essentials: Ridge, Lasso & Elastic Net"