



STA 235 - Prediction II: K-nearest neighbors

Spring 2021

McCombs School of Business, UT Austin


Prediction tasks

- We have seen the main issue with **bias vs variance trade-off**
- Beyond regression, **what methods can we use for prediction?**


K-nearest neighbor

KNN as a classification problem

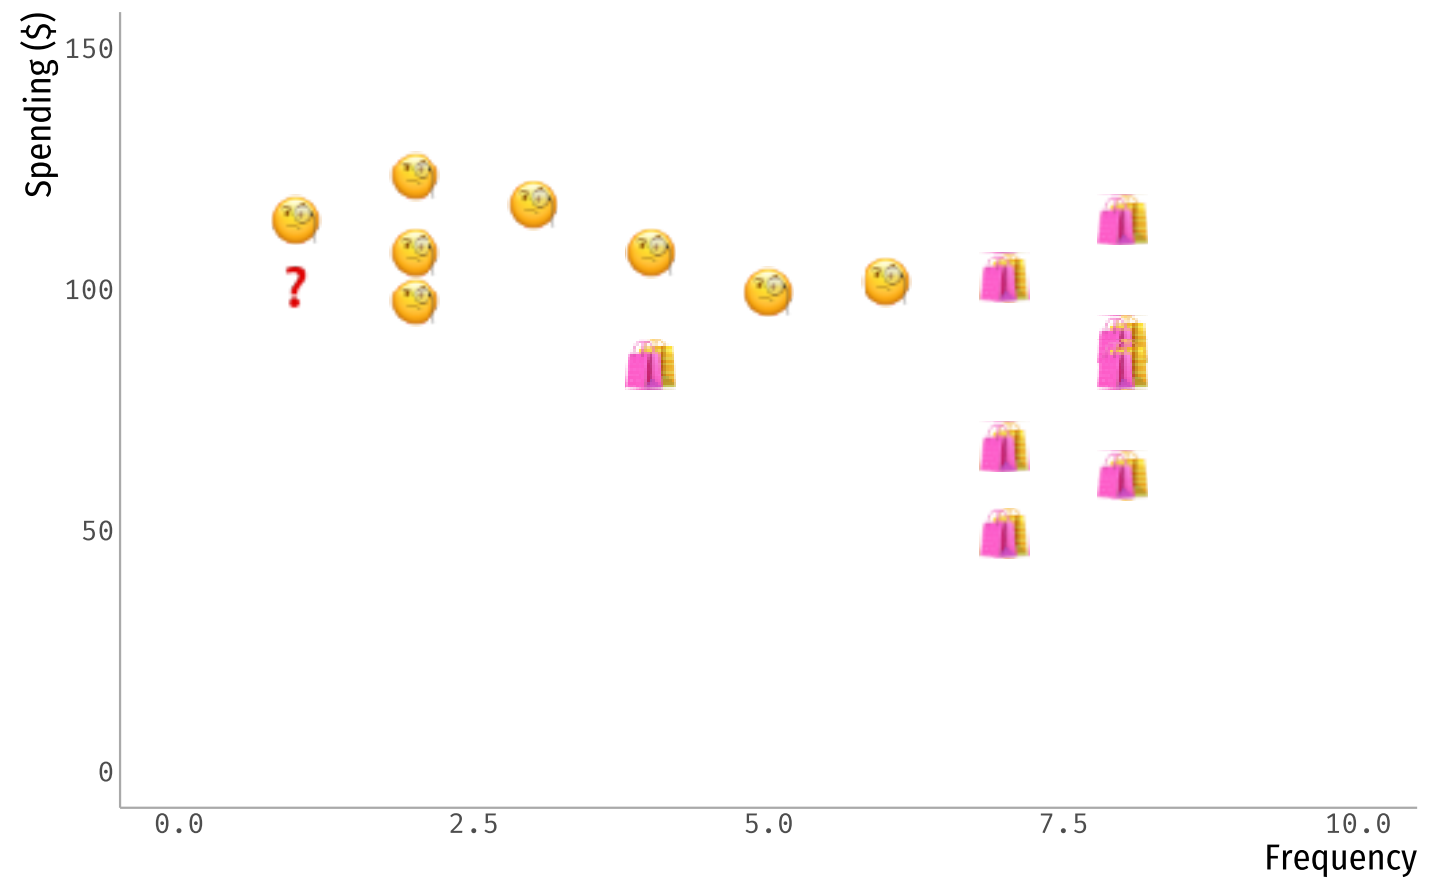
- Again: Window shoppers vs high rollers




How would you classify this unit?



How would you classify this unit?



But what about this one?



K-nearest neighbor classifier

- One of the **simplest classifications methods**

- 1) Choose a **distance measure** (e.g. euclidian).
- 2) Choose a **number of neighbors**, K (*Note: Choose an odd number!*).
- 3) **Calculate the distance** between data and other points.
- 4) Calculate the **rate for each class** according to K : $Pr(Y = j|X = x_0) = \frac{1}{K} \sum_{i \in N_0} \mathbf{I}(y_i = j)$.
- 5) **Assign the majority class**.

KNN with $K = 1$

Classifier: High-roller

KNN with $K = 3$

Classifier: High-roller

KNN with $K = 9$

Classifier: Window-shopper

Poll time!

**A lower number of neighbors K
yields...**

KNN Classifier in R?

```
library(caret)

d <- read.csv("https://raw.githubusercontent.com/maibennett/sta235/main/exampleSite/content/Classes",
head(d)
```

```
##      freq female spend type
## 1     10       1    59   WS
## 2      7       1    71   WS
## 3      6       1    79   WS
## 4      3       0    97   HR
## 5      9       1    52   WS
## 6     10       1    56   WS
```

KNN Classifier in R?

```
library(caret)

d <- read.csv("https://raw.githubusercontent.com/jjallaire/ml/master/data/sonar.csv")

set.seed(100)

n <- nrow(d)

train.row <- sample(1:n, 0.8*n)

test.data <- d[-train.row,]
train.data <- d[train.row,]

knn <- train(
  type ~., data = train.data,
  method = "knn",
  trControl = trainControl("cv", number = 10),
  preProcess = c("center", "scale"),
  tuneLength = 15
)
```

- Again, we'll be using the `caret` package.

KNN Classifier in R?

```
library(caret)

d <- read.csv("https://raw.githubusercontent.com/...")
set.seed(100)

n <- nrow(d)

train.row <- sample(1:n, 0.8*n)

test.data <- d[-train.row,]
train.data <- d[train.row,]

knn <- train(
  type ~., data = train.data,
  method = "knn",
  trControl = trainControl("cv", number = 10),
  preProcess = c("center", "scale"),
  tuneLength = 15
)
```

- Again, we'll be using the `caret` package.
- Create a **training** and **testing** dataset.

KNN Classifier in R?

```
library(caret)

d <- read.csv("https://raw.githubusercontent.com/josiah-davis/100-questions-in-100-days-of-r/master/data/100_questions.csv")

set.seed(100)

n <- nrow(d)

train.row <- sample(1:n, 0.8*n)

test.data <- d[-train.row,]
train.data <- d[train.row,]

knn <- train(
  type ~., data = train.data,
  method = "knn",
  trControl = trainControl("cv", number = 10),
  preProcess = c("center", "scale"),
  tuneLength = 15
)
```

- Again, we'll be using the `caret` package.
- Create a **training** and **testing** dataset.
- Use the method `knn` on a factor variable (i.e. classification)

KNN Classifier in R?

```
library(caret)

d <- read.csv("https://raw.githubusercontent.com/...")
set.seed(100)

n <- nrow(d)

train.row <- sample(1:n, 0.8*n)

test.data <- d[-train.row,]
train.data <- d[train.row,]

knn <- train(
  type ~., data = train.data,
  method = "knn",
  trControl = trainControl("cv", number = 10),
  preProcess = c("center", "scale"),
  tuneLength = 15
)
```

- Again, we'll be using the **caret** package.
- Create a **training** and **testing** dataset.
- Use the method **knn** on a factor variable (i.e. classification)
- We also **pre-process** the data. Why?

KNN Classifier in R?

```
library(caret)

d <- read.csv("https://raw.githubusercontent.com/...")
set.seed(100)

n <- nrow(d)

train.row <- sample(1:n, 0.8*n)


test.data <- d[-train.row,]
train.data <- d[train.row,]

knn <- train(
  type ~., data = train.data,
  method = "knn",
  trControl = trainControl("cv", number = 10),
  preProcess = c("center", "scale"),
  tuneLength = 15
)
```

- Again, we'll be using the `caret` package.
- Create a **training** and **testing** dataset.
- Use the method `knn` on a factor variable (i.e. classification)
- We also **pre-process** the data. Why?
- `tuneLength` is the level of granularity for searching K .


How many neighbors?

We can see the optimal K using `bestTune` parameter.



Poll time

Which K would you choose?



How accurate is this?

- For **classification** problems, we care about *false positive* and *false negative*.
- Say 1: window-shoppers and 2: high-rollers.

```
pred.type <- knn %>% predict(test.data)
table(pred.type, test.data$type)
```

```
##
## pred.type HR WS
##          HR 72 28
##          WS 17 83
```

Poll time

In a table like this, where would you like to see most of the observations?

```
pred.type <- knn %>% predict(test.data)
table(pred.type, test.data$type)
```

```
##
## pred.type HR WS
##          HR 73 28
##          WS 16 83
```


How accurate is this?

- For **classification** problems, we care about *false positive* and *false negative*.
- Say 1: window-shoppers and 2: high-rollers.

```
pred.type <- knn %>% predict(test.data)
table(pred.type, test.data$type)
```

```
##
## pred.type HR WS
##          HR 73 29
##          WS 16 82
```

```
mean(pred.type == test.data$type)
```

```
## [1] 0.775
```

KNN for regression

- We can also use KNN for **continuous outcomes**
- **Similar** to the KNN classifier, but now we will take the *average of the K-neighbors* for prediction:

$$\hat{f}(x_0) = \frac{1}{K} \sum_{i \in N_0} y_i$$

KNN Regression in R?

```
library(caret)

d <- read.csv("https://raw.githubusercontent.com/...")

set.seed(100)

n <- nrow(d)

train.row <- sample(1:n, 0.8*n)

test.data <- d[-train.row,]
train.data <- d[train.row,]


knnr <- train(
  spend ~. - type, data = train.data,
  method = "knn",
  trControl = trainControl("cv", number = 10),
  preProcess = c("center", "scale"),
  tuneLength = 50
)
```

Same as before!

... but with a continuous variable


Choose optimal K

We get the optimal K the same way, using `knnr$bestTune`



Poll time

Which K would you choose?



Takeaway points



- KNN is a simple, nonparametric way to do prediction for both **categorical** and **continuous** outcomes.
- Be sure to **check your accuracy/error metric** depending on your outcome.
- **Pre-processing** can play an important role!

Plot your data and results

Next class

- Other **prediction methods**:

Decision trees!



References

- James, G. et al. (2013). "Introduction to Statistical Learning with Applications in R". *Springer. Chapter 2, Chapter 3.*
- STDHA. (2018). "KNN: K-Nearest Neighbors Essentials"